

Heart Disease Classification Problem

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
```

```
In [5]: df = pd.read_csv("heart-disease.csv")
df
```

```
Out[5]:
```


	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

303 rows × 14 columns

```
In [6]: df.target.value_counts(normalize=True) * 100
```

```
Out[6]: 1 ... 54.455446
0 ... 45.544554
Name: target, dtype: float64
```

```
In [7]: df.target.value_counts().plot(kind="bar",color=["salmon", "lightblue"]);
```




```
In [6]: pd.crosstab(df.target,df.cp)
```

```
Out[6]:
```

	cp	0	1	2	3
target					
0		104	9	18	7
1		39	41	69	16

```
In [7]: plt.figure(figsize=(12, 12))
sns.heatmap(df.corr(),
            annot=True,
            linewidths=0.5,
            fmt= ".2f",
            );
```



```
In [8]: x = df.drop("target", axis= 1)
y = df.target
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

```
In [9]: from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
models = {"Random Forest":RandomForestClassifier(),"Logistic Regression":LogisticRegression()
```

```
In [10]: def fit_and_score(models,x_train, x_test, y_train, y_test):
models_scores = {}
for name,model in models.items():
model.fit(x_train, y_train)
models_scores[name] = model.score(x_test, y_test)
return models_scores
models_scores = fit_and_score(models,x_train, x_test, y_train, y_test)
models_scores
```


C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

Out[10]: {'Random Forest': 0.819672131147541, 'Logistic Regression': 0.8688524590163934}

In [11]: models_df = pd.DataFrame(models_scores, index=["accuracy"])
models_df.T.plot(kind="bar")

Out[11]: <AxesSubplot:>



In [12]: lr_grid = {"C":np.logspace(-4,4,20),"solver":["liblinear"]}
lr_gs = GridSearchCV(LogisticRegression(),lr_grid,cv = 5,verbose=True)
lr_gs.fit(x_train, y_train)

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
Out[12]: GridSearchCV(cv=5, estimator=LogisticRegression(),
.....:      param_grid={'C': array([1.00000000e-04, 2.63665090e-04, 6.95192796e-04, 1.
.....:      83298071e-03,
.....:      4.83293024e-03, 1.27427499e-02, 3.35981829e-02, 8.85866790e-02,
.....:      2.33572147e-01, 6.15848211e-01, 1.62377674e+00, 4.28133240e+00,
.....:      1.12883789e+01, 2.97635144e+01, 7.84759970e+01, 2.06913808e+02,
.....:      5.45559478e+02, 1.43844989e+03, 3.79269019e+03, 1.00000000e+04]),
.....:      'solver': ['liblinear']}),
.....:      verbose=True)
```

```
In [13]: lr_gs.best_params_
```

```
Out[13]: {'C': 0.08858667904100823, 'solver': 'liblinear'}
```

```
In [14]: lr_gs.score(x_test, y_test)
```

```
Out[14]: 0.8360655737704918
```

```
In [15]: models_scores
```

```
Out[15]: {'Random Forest': 0.819672131147541, 'Logistic Regression': 0.8688524590163934}
```


```
In [16]: model = models["Random Forest"]
```

```
In [18]: from sklearn.metrics import plot_roc_curve, confusion_matrix
```

```
# Plot ROC curve and calculate AUC metric
plot_roc_curve(model, x_test, y_test)
plt.savefig(fname="png.png")
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDisplay.from_estimator`.

warnings.warn(msg, category=FutureWarning)



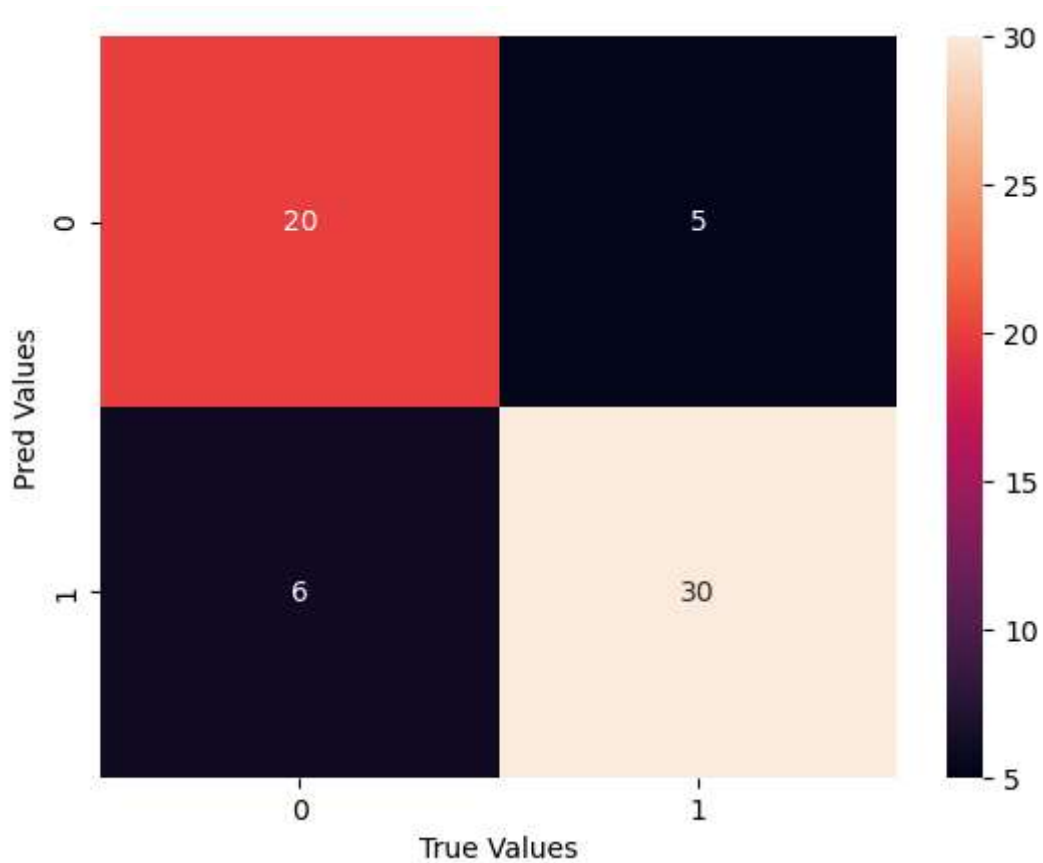
```
In [23]: y_preds = model.predict(x_test)
         y_preds
```

```
Out[23]: array([0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1,
        .....: 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1,
        .....: 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1], dtype=int64)
```

```
In [27]: confusion_matrix(y_test, y_preds)
```

```
Out[27]: array([[20,  5],
        .....: [ 6, 30]], dtype=int64)
```

```
In [32]: fig, ax = plt.subplots()
         ax = sns.heatmap(confusion_matrix(y_test, y_preds),
         .....:          annot = True,
         .....:          )
         plt.xlabel("True Values")
         plt.ylabel("Pred Values")
         fig.savefig("Confusion Matrix")
```



```
In [36]: from sklearn.model_selection import cross_val_score

cv_acc = np.mean(cross_val_score(model,
                                  x,
                                  y,
                                  cv=5, # 5-fold cross-validation
                                  scoring="accuracy"))

cv_acc
```

Out[36]: 0.8182513661202186

```
In [38]: cv_pre = np.mean(cross_val_score(model,
                                             x,
                                             y,
                                             cv=5, # 5-fold cross-validation
                                             scoring="precision"))

cv_pre
```

Out[38]: 0.8284248053288301

```
In [41]: cv_recall = np.mean(cross_val_score(model,
                                                x,
                                                y,
                                                cv=5, # 5-fold cross-validation
                                                scoring="recall"))

cv_recall
```

Out[41]: 0.8727272727272727

```
In [43]: cv_f1 = np.mean(cross_val_score(model,  
                                   x,  
                                   y,  
                                   cv=5, # 5-fold cross-validation  
                                   scoring="f1"))  
  
cv_f1
```

Out[43]: 0.8442794468492067

```
In [55]: # Visualizing cross-validated metrics  
cv_metrics = pd.DataFrame({"Accuracy": cv_acc,  
                           "Precision": cv_pre,  
                           "Recall": cv_recall,  
                           "F1": cv_f1}, index=[0])  
cv_metrics.T.plot(kind="bar", title="Cross-Validated Metrics", legend=False);  
plt.savefig("Cross-Validated Metrics")
```

